

AUTOML: UNA BREVE INTRODUCCIÓN


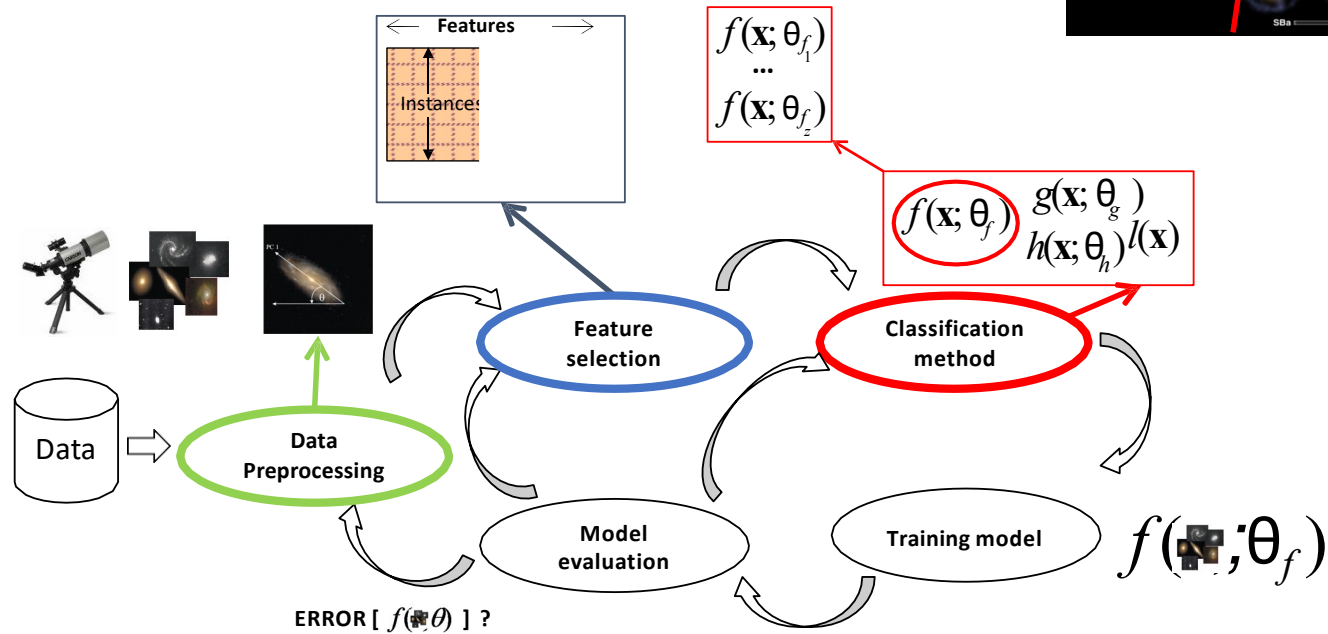
Minería de Datos: Aspectos Avanzados

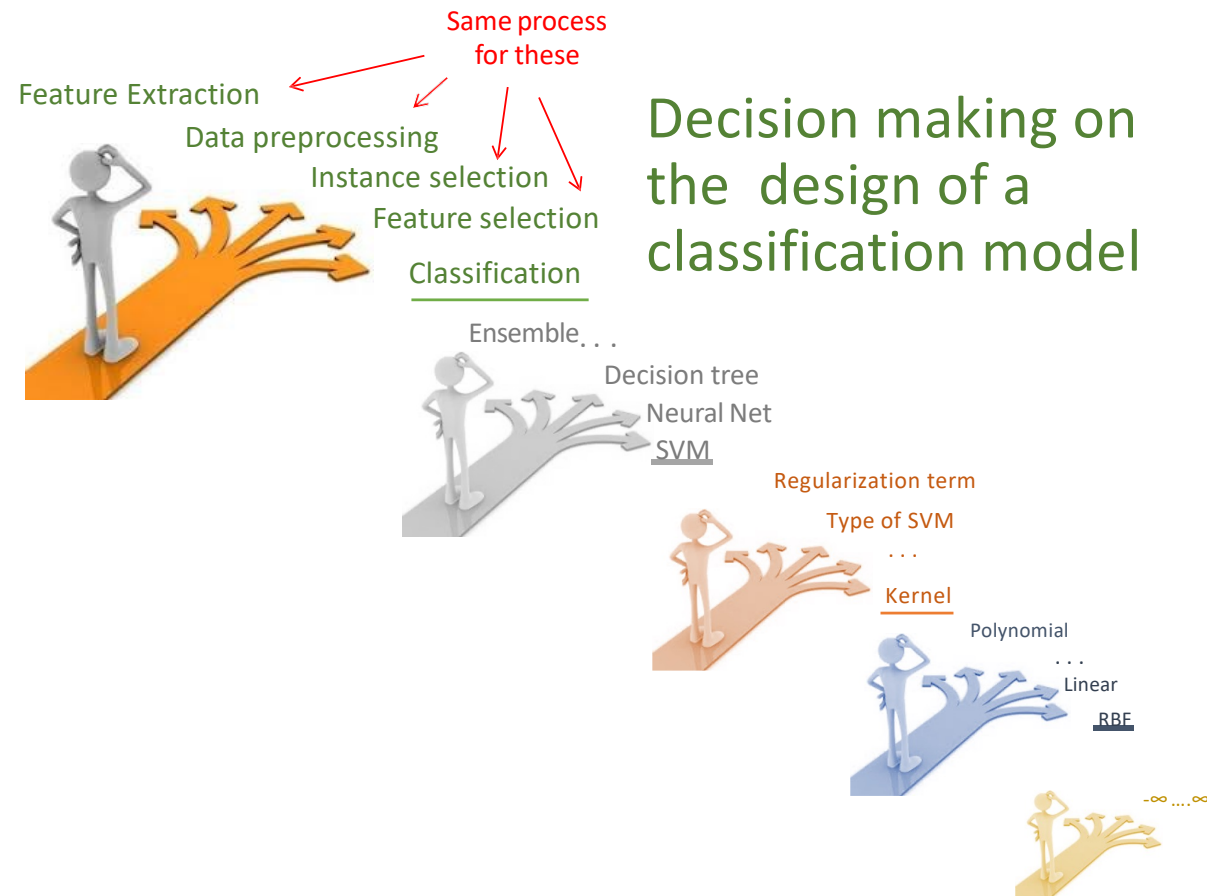
Salvador García

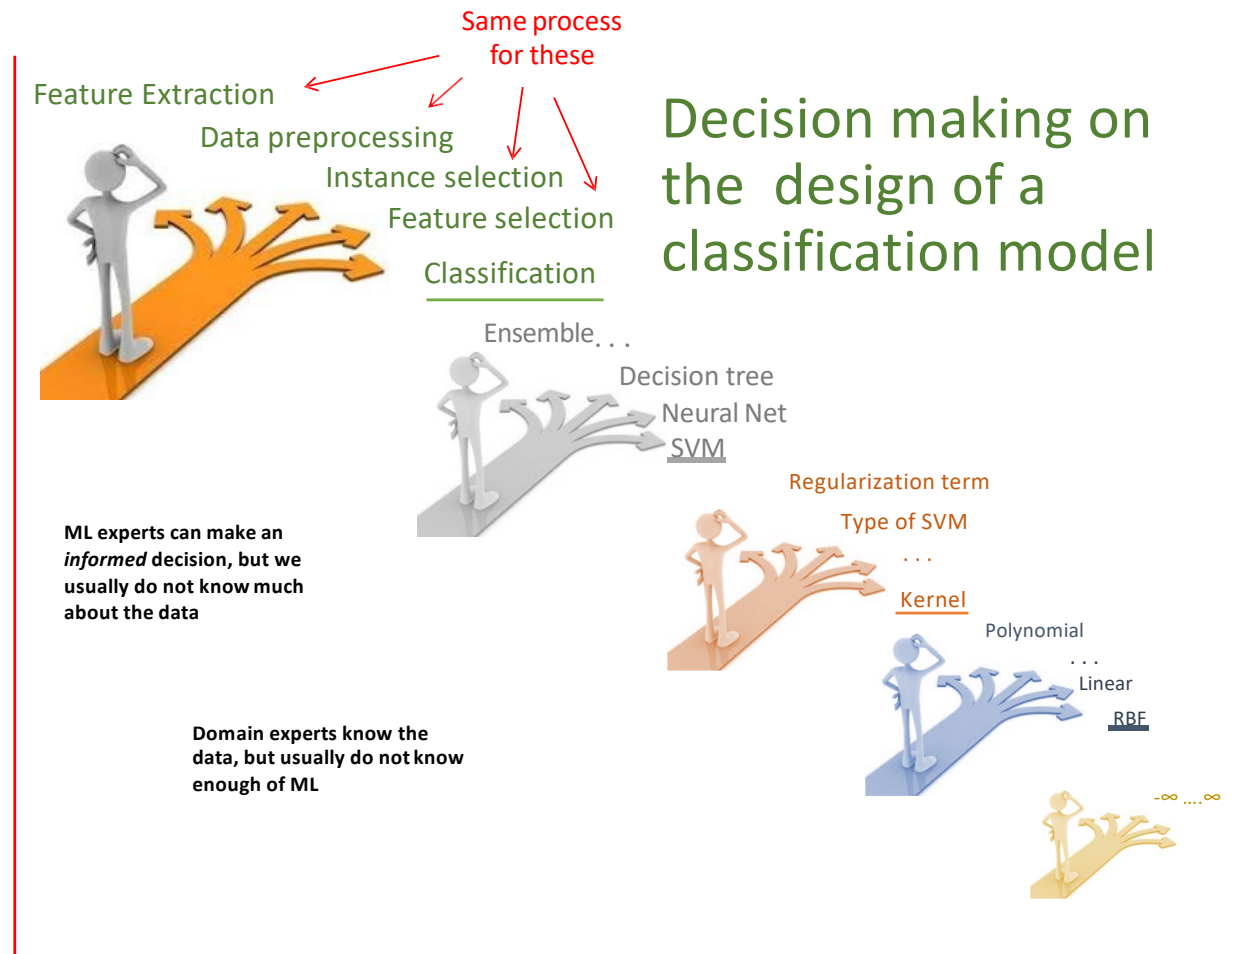
salvagl@decsai.ugr.es

ML vs AutoML

Classification: Typical design process

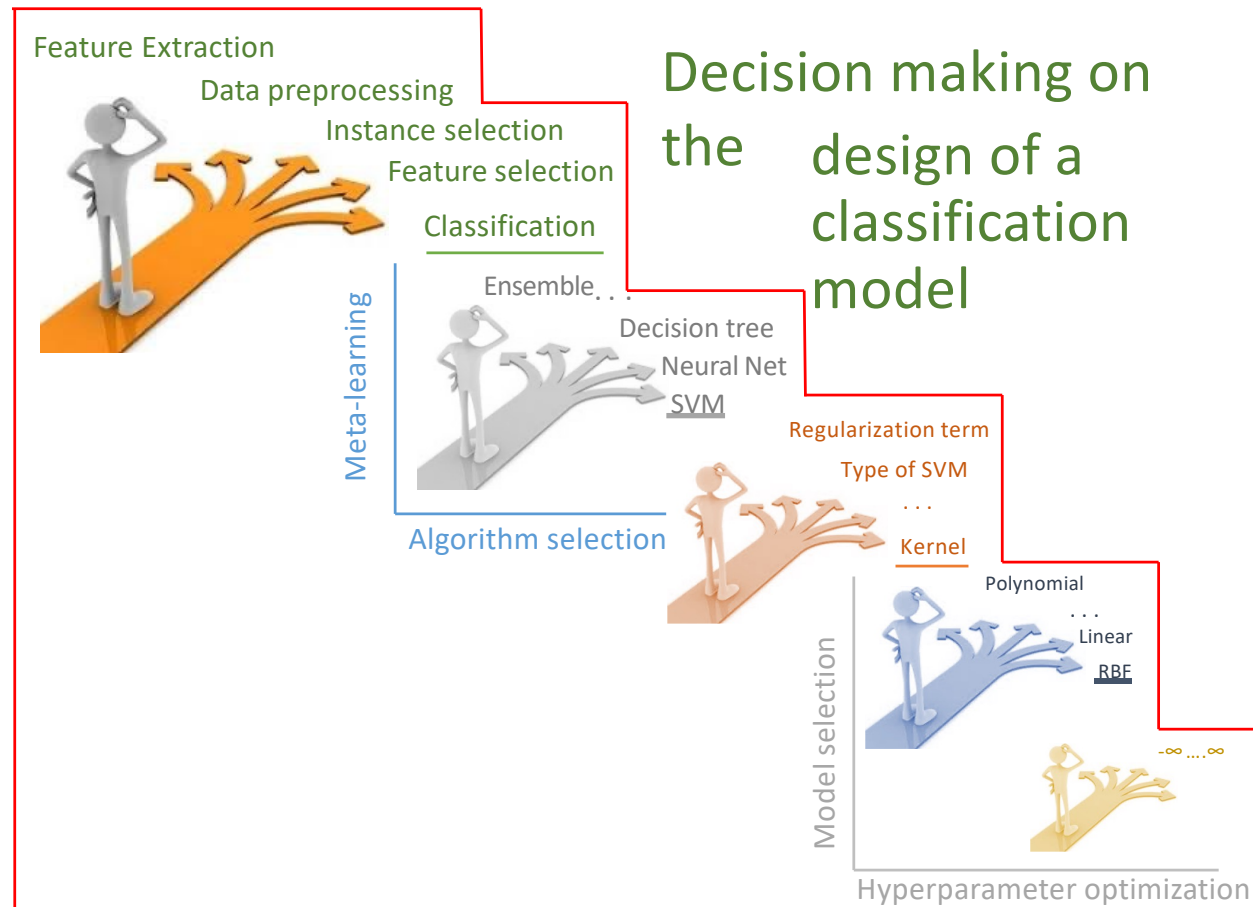






Some issues with the cycle of design


- Commonly, the above issues are fixed manually, relying on:
 - Domain expert's knowledge
 - Machine learning specialists' knowledge
 - Trial and error
- The design/development of a pattern classification system relies on the knowledge and biases of humans, which may be risky, expensive and time consuming
- ***Automated solutions are available but only for particular processes (e.g., either feature selection, or classifier selection but not both)***




AutoML

- Automatic Machine Learning*
 - Research area that targets progressive automation of machine learning
 - Field of research focusing on the development of autonomous methods for solving a variety of machine learning problems

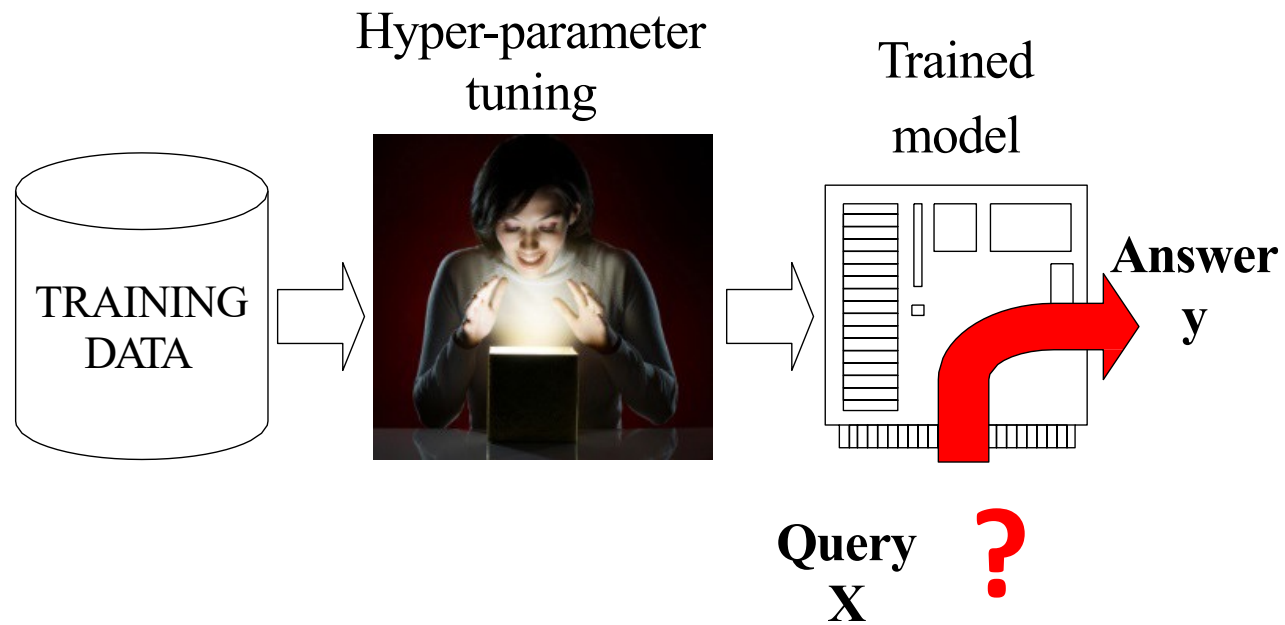
* We will focus on supervised learning



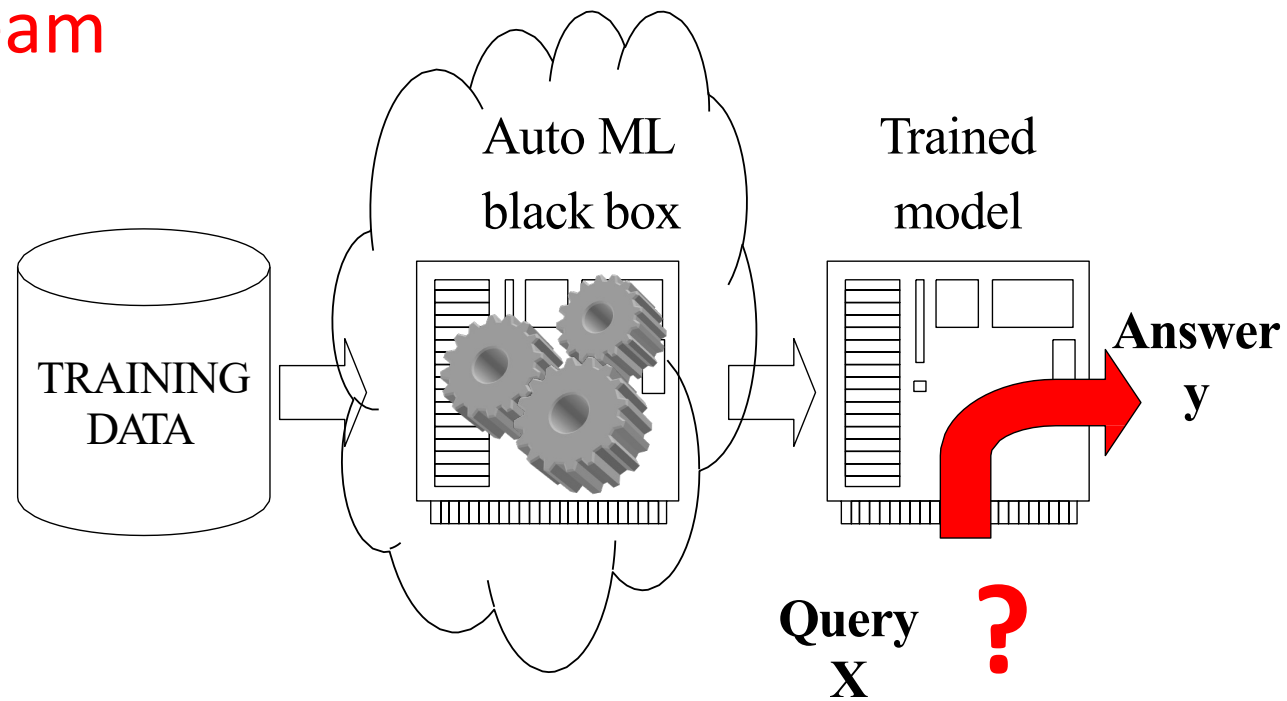
ML: The Problem



ML: The REALITY



AutoML: The Dream



Why AutoML?

- Relevance
 - Large amounts of data readily available everywhere
 - Lack of domain and/or ML experts who can advise/supervise the development of ML-based systems
 - Need to tune ML models
- Complexity
 - Huge and heterogeneous search space
 - Overfitting

AutoML notions

Informal, but intuitive definition

- *AutoML is the task of finding the model (f) that better generalizes in any possible dataset (T) with the less possible human intervention*
 - f can be the composition of multiple functions that may transform the input space, subsampling data, combining multiple predictors, etc.
- where each of these models could be formed in turn by several other functions/models.

An inclusive notion of AutoML

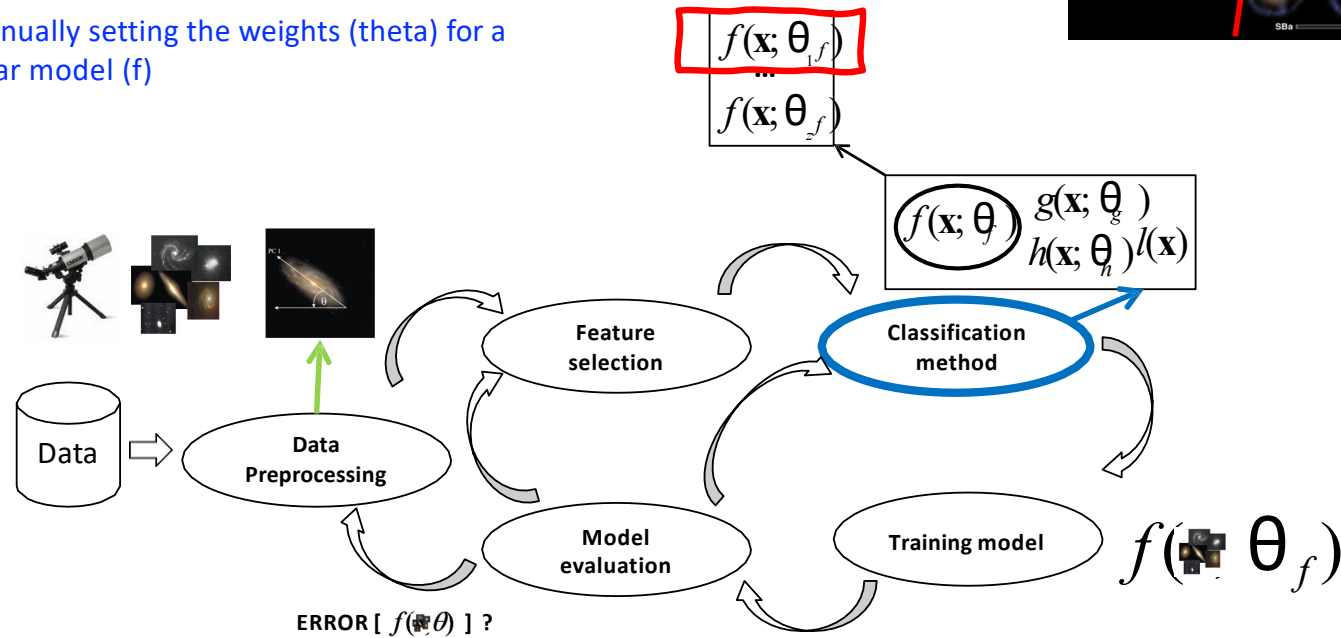

- Three level categorization
 - Alpha-level - Search of estimators
 - Beta-level – Search of learning algorithms
 - Gamma-level – Search for meta-learning algorithms

Level	Input	Output	Examples	Encoded by
α -level	sample/example (e.g. an image)	prediction of label (e.g. 'dog' or 'cat')	heuristically hard-coded classifier or already trained classifier	parameters, hyperparameters (if any) and meta- parameters (if any)
β -level	task/dataset (e.g. MNIST [20], CIFAR-10 [19])	α -level algorithm	learning algorithms (e.g. SVM [6], CNN [20]); HPO algorithms (e.g. grid search cross-validation, SMAC [1], NAS [34])	hyperparameters and meta-parameters (if any)
γ -level	meta-dataset (e.g. OpenML [29])	β -level algorithm	meta-learning algorithms (e.g. meta-learning part in Auto-sklearn [11])	meta-parameters

Liu Zhengying, Zhen Xu, Meysam Madadi, Julio Jacques Junior, Sergio Escalera, Shangeth Rajaa, and Isabelle Guyon. **Overview and unifying conceptualization of automated machine learning**. In Proc. of Automating Data Science Workshop @ECML-PKDD, 2019.


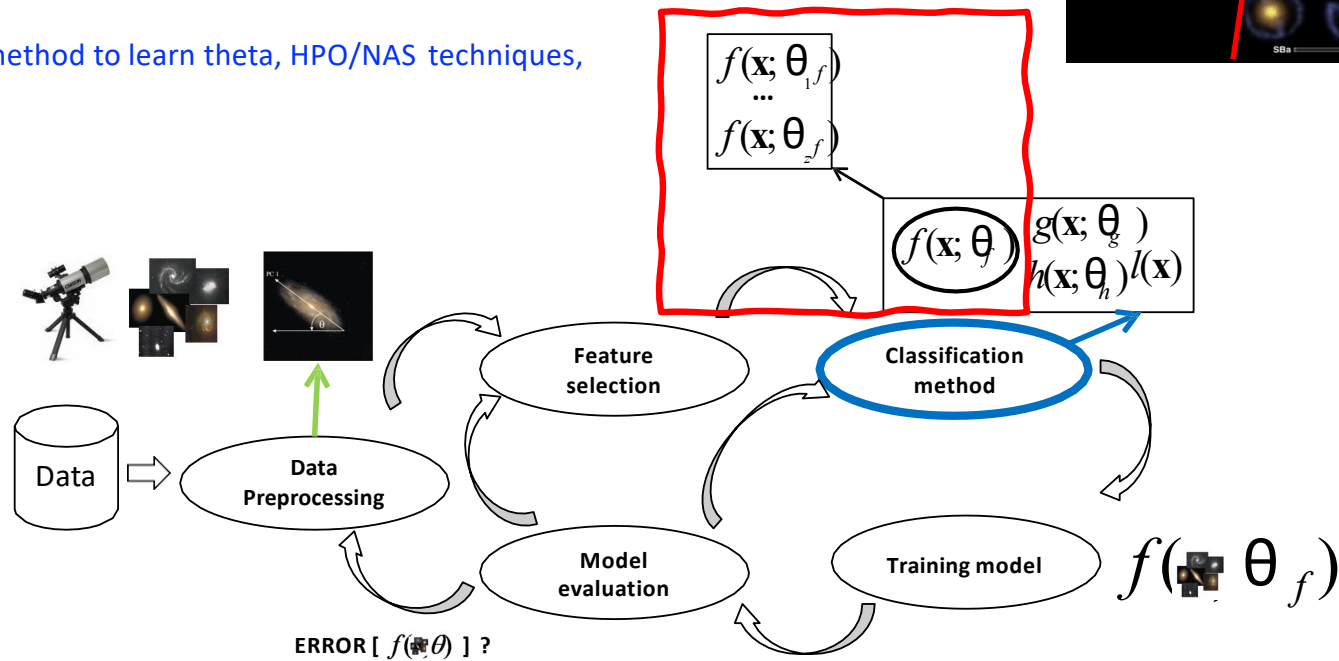
Alpha level

e.g., Manually setting the weights (theta) for a particular model (f)



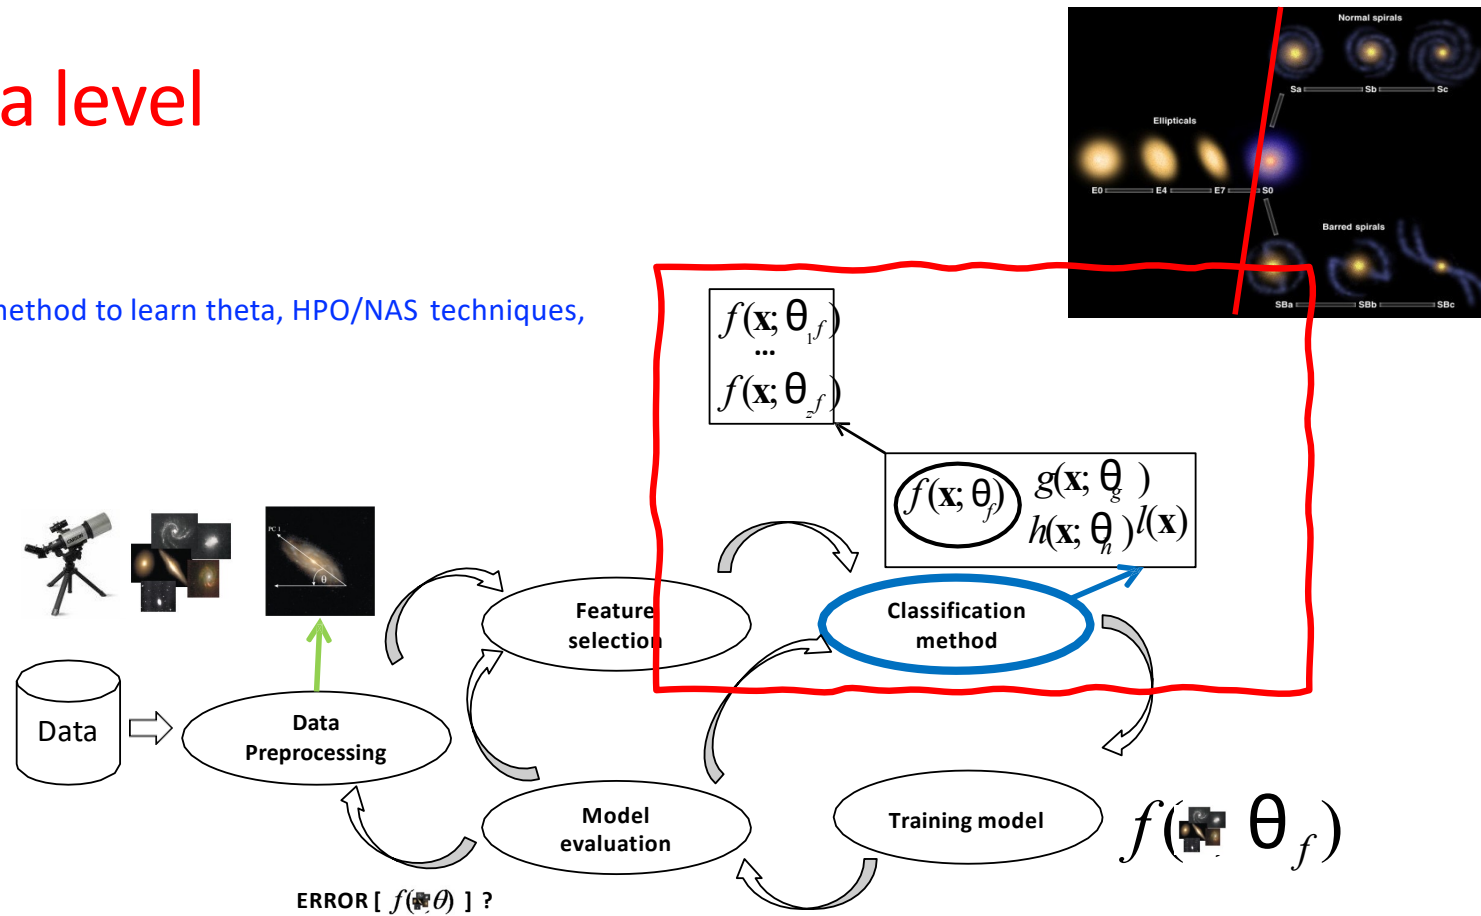
Beta level

e.g., A method to learn theta, HPO/NAS techniques,



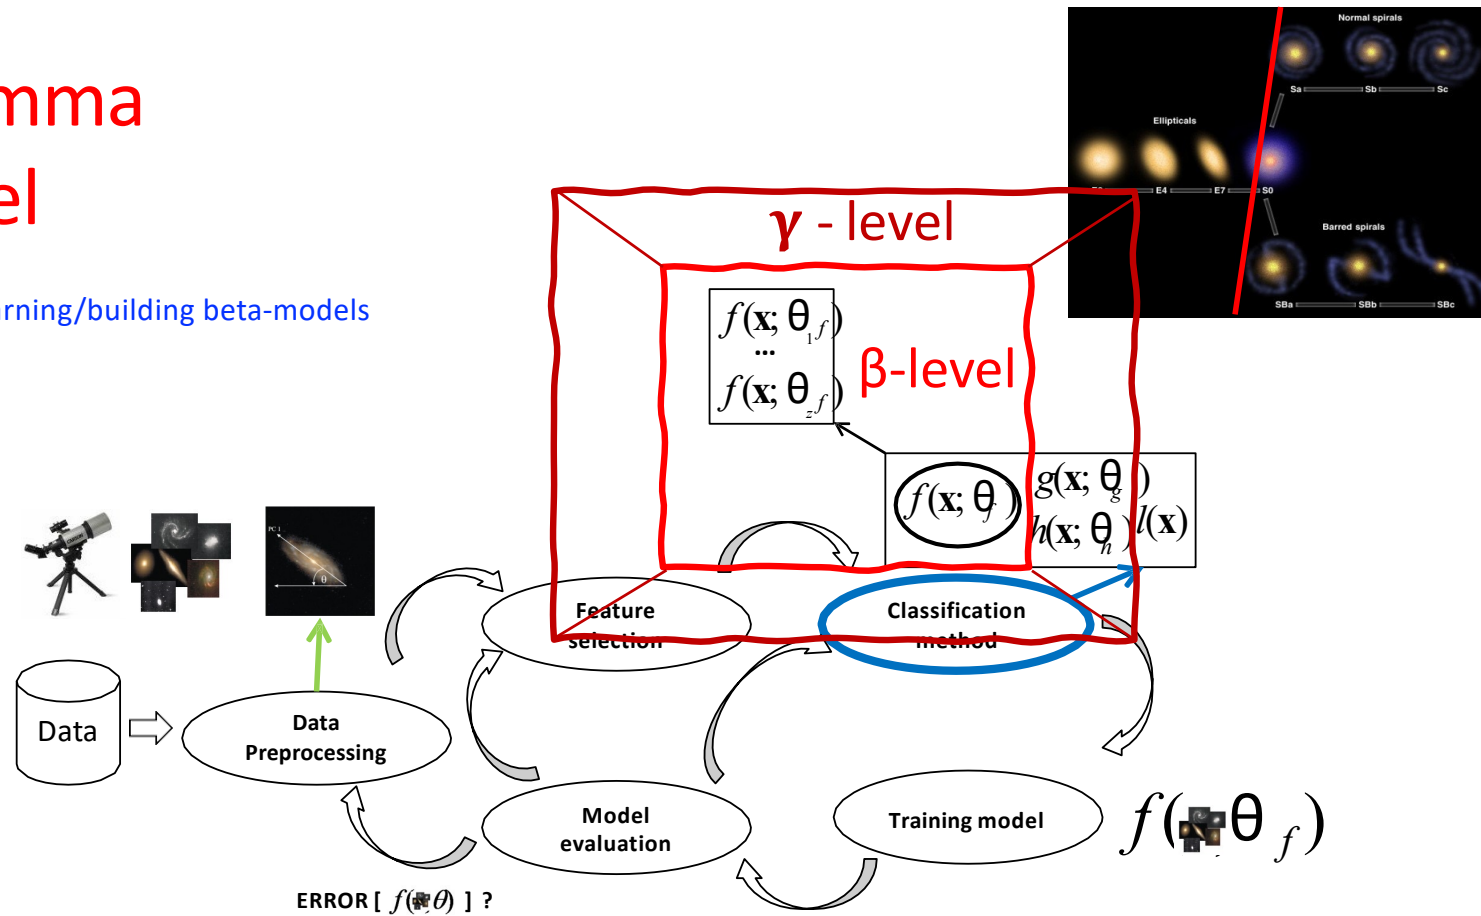
Beta level

e.g., A method to learn theta, HPO/NAS techniques,



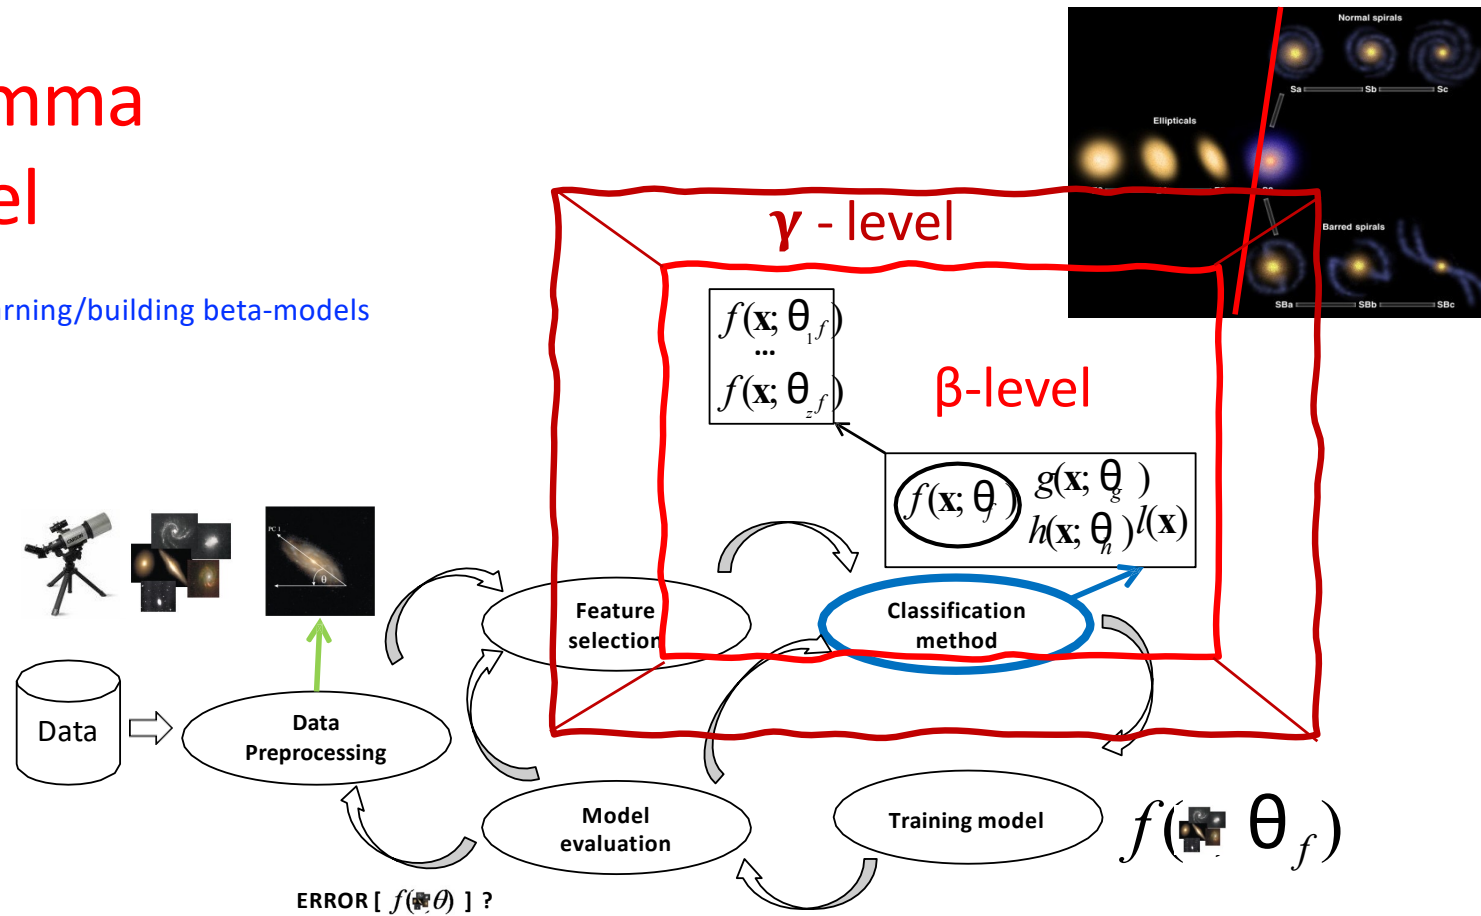
Gamma level

e.g., Learning/building beta-models



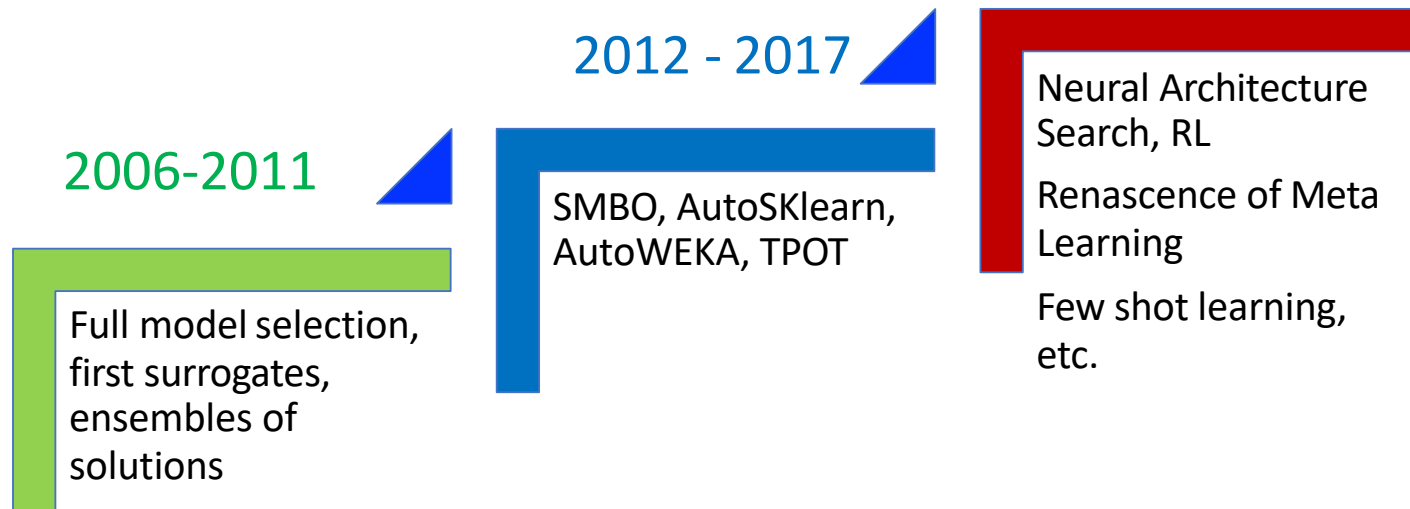
Gamma level

e.g., Learning/building beta-models



The three waves of AutoML

Three waves of AutoML





First wave (2006 – 2011)


- Meta learning (early efforts)
 - Algorithm selection, classifier recommendation
- Full model (pipeline) selection vs HPO
- Ensembles of partial solutions
- Surrogate models

PSMS : PSO for full model selection

- Codification of solutions as real valued vectors





- Using PSO straightforwardly




PSMS : PSO for full model selection

- Codification of solutions as real valued vectors



- Using PSO straightforwardly



Main findings from the early years

- Overfitting avoidance mechanisms
- Heterogeneous codification of solutions
- “Straightforward” global optimization approaches
- Data subsampling for efficient estimation of the objective function
- Assembling solutions
- Template-based initialization


Second wave (2012-2017)

- Sequential Model Based Optimization
- Formulation of the CASH (Combined Algorithm Selection and Hyperparameter optimization) problem (quite similar to FMS)
- Multi objective approaches to FMS
- Meta learning + Optimization

Bayesian optimization

- A global optimization procedure, based on two functions:
 - Surrogate,
 - Gaussian processes, random forest, etc
 - Acquisition functions

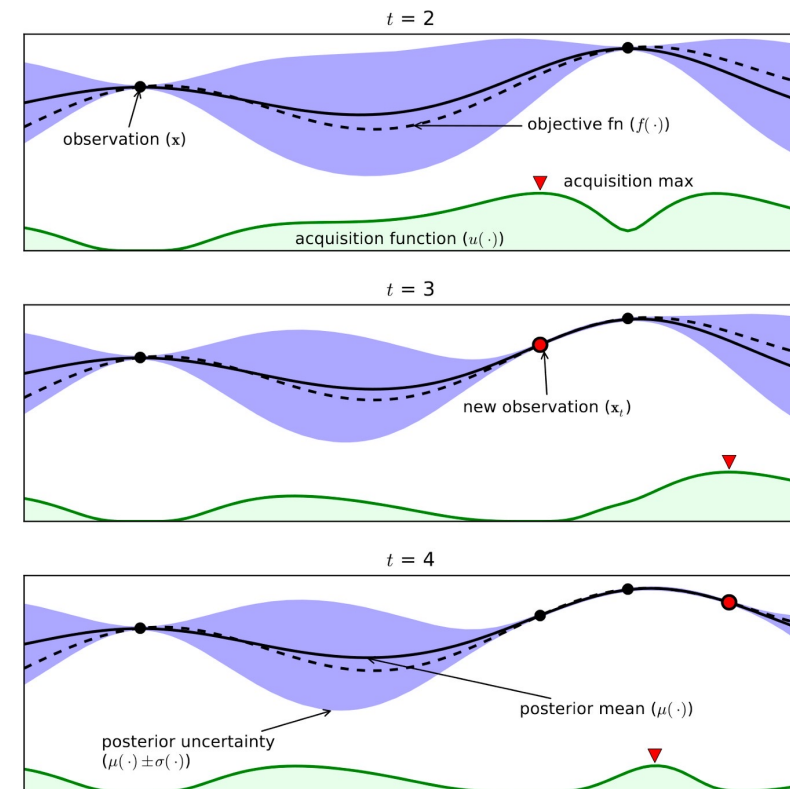
$$\mathbb{E}[\mathbb{I}(\lambda)] = \mathbb{E}[\max(f_{\min} - y, 0)]$$



AutoWeka

- Introduced the definition of CASH, and approached the problem with Bayesian Optimization / Sequential Model-Based Optimization (SMBO)
 - BO/SMBO: sequential design strategy for global optimization of black-box functions that doesn't require derivatives.

<https://www.cs.ubc.ca/labs/beta/Projects/autoweka/>




C. Thornton, F. Hutter, H. Hoos, and K. Leyton-Brown. Auto-WEKA: combined selection and hyperparameter optimization of classification algorithms. In Proc. KDD'13, pages 847–855, 2013.

Multi objective FMS



- Selecting models that optimize more than a single criterion combinations include:
 - Bias-variance
 - Performance-time
 - Performance-complexity

$$\begin{array}{ll}\text{minimize} & \mathbf{f}(\mathbf{x}) = [f_1(\mathbf{x}), \dots, f_l(\mathbf{x})]^T \\ \text{subject to} & \mathbf{x} \in \mathcal{X}\end{array}$$



Multi objective FMS

- Selecting models that optimize more than a single criterion



Surrogate models
were used here!

Alejandro Rosales-Pérez, Jesus A. Gonzalez, Carlos A. Coello Coello, Hugo Jair Escalante, Carlos A. Reyes García: **Multi-objective model type selection**. Neurocomputing 146: 83-94 (2014)


Alejandro Rosales-Pérez a,n, Jesus A. Gonzalez a, Carlos A. Coello Coello b, Hugo Jair Escalante a, Carlos A. Reyes-Garcia. **Surrogate-assisted multi-objective model selection for support vector machines**. Neurocomputing 150 (2015) 163–172, 2015

Multi fidelity approaches


- For AutoML solutions based on black box optimization, the evaluation of a candidate solution (model) is computationally expensive, and usually the more models are evaluated the better the performance of the AutoML methodology
- Methods aiming to make budget-constraint approximations of the real performance have been proposed

Multi fidelity approaches


- Learning curve-based performance:
 - Build and try to predict a learning curve (e.g., model performance vs. dataset size or model complexity)
 - Discard the model whenever the predictive curve is not promising



(a) SMAC on k-means CIFAR-10



(b) TPE on k-means CIFAR-10




(c) SMAC on MNIST

Tobias Domhan, Jost Tobias Springenberg, Frank Hutter **Speeding up Automatic Hyperparameter Optimization of Deep Neural Networks by Extrapolation of Learning Curves.** Proc. IJCAI 2015

Multi fidelity approaches

- **Successive halving (SH):** Given a budget (time, other resources), evaluate all algorithms for that budget; then remove the half worst models from consideration; duplicate the budget and repeat until a single model remains.
- **Hyperband:** Generate random configurations of different budgets. Then using SH as subroutine



Algorithm 1: HYPERBAND algorithm for hyperparameter optimization.

```

input      :  $R, \eta$  (default  $\eta = 3$ )
initialization:  $s_{\max} = \lfloor \log_{\eta}(R) \rfloor, B = (s_{\max} + 1)R$ 
1 for  $s \in \{s_{\max}, s_{\max} - 1, \dots, 0\}$  do
2    $n = \lceil \frac{B}{R} \frac{\eta^s}{(s+1)} \rceil, \quad r = R\eta^{-s}$ 
   // begin SUCCESSIVEHALVING with  $(n, r)$  inner loop
3    $T = \text{get\_hyperparameter\_configuration}(n)$ 
4   for  $i \in \{0, \dots, s\}$  do
5      $n_i = \lfloor n\eta^{-i} \rfloor$ 
6      $r_i = r\eta^i$ 
7      $L = \{\text{run\_then\_return\_val\_loss}(t, r_i) : t \in T\}$ 
8      $T = \text{top\_k}(T, L, \lfloor n_i/\eta \rfloor)$ 
9   end
10 end
11 return Configuration with the smallest intermediate loss seen so far.

```

K. Jamieson and Talwalkar. A. Non-stochastic Best Arm Identification and Hyperparameter Optimization. 2016

L. Li et al. Hyperband: A Novel Bandit-Based Approach to Hyperparameter Optimization. JMLR 2018

Meta-learning: learning to learn


- Evaluates and compares the application of learning algorithms on (many) previous tasks/domains to suggest learning algorithms (combinations, rankings) for new tasks
- Focuses on the relation between tasks/domains and learning algorithms
- Accumulating experience on the performance of multiple applications of learning methods

Brazdil P., Giraud-Carrier C., Soares C., Vilalta R. **Metalearning: Applications to Data Mining**. Springer Verlag. ISBN: 978-3-540-73262-4, 2008.

Brazdil P., Vilalta R., Giraud-Carrier C., Soares C.. **Metalearning**. Encyclopedia of Machine learning. Springer, 2010.

AutoSKlearn

- Meta-learning + SMBO
- Trained a meta-learner on OpenML, uses it to warmstart the optimization process
- Ensembles



Matthias Feurer, Aaron Klein, Katharina Eggenberger, Jost Tobias Springenberg, Manuel Blum, and Frank Hutter. Auto-sklearn: Efficient and Robust Automated Machine Learning, pages 113–134. Springer International Publishing, Cham, 2019.

Main developments from the second wave

- Surrogate models
- Multi objective AutoML
- Bayesian optimization is now a “standard”
- Multi fidelity approaches
- Resurgence of meta-learning
- ...

Three waves of AutoML

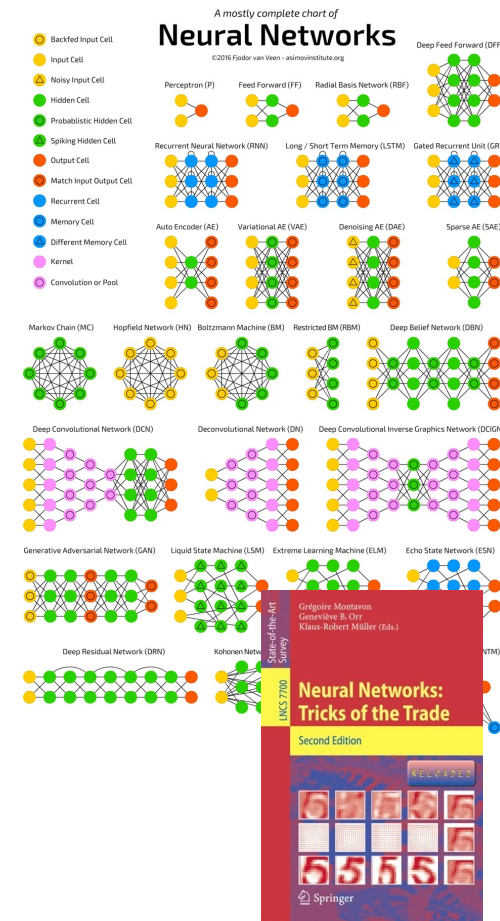


Third wave (2018 - and on ...)

- Neural architecture search
- Reinforcement learning based AutoML
- Reinassence of evolutionary algorithms for AutoML
- AutoML from raw data!


Neural architecture search

- It is well known that the success of deep learning (DL) solutions relies heavily on the design choices and hyperparameters
- DL is everywhere, hence, AutoML methods targeting DL can have a huge impact



Neural architecture search

- NAS is the process of automating architecture design in the context of DL models (share same complications as AutoML in supervised learning)
- Components of a NAS method:
 - Search space
 - Search strategy
 - Performance estimation strategy



Neural architecture search

- Search strategy:
 - Evolutionary algorithms
 - Bayesian optimization
 - Reinforcement learning
 - Gradient based
- Performance estimation:
 - Multi fidelity approaches
 - One shot architecture search

Latest developments

- Neural architecture search is driving new ways for efficient AutoML
- RL-based AutoML
- Meta learning resurgence
- Complex search spaces that require of new optimization methodologies

Final remarks

AutoML / Full model selection

- Pros
 - The job of the data analyst is considerably reduced
 - Neither knowledge on the application domain nor on machine learning is required
 - Different methods for preprocessing, feature selection and classification are considered
 - It can be used in any classification problem
- Cons
 - It is real function + combinatorial optimization problem
 - Computationally expensive
 - Risk of overfitting

NLF theorem!



... and full models for all

- **Short-term goal:** to provide data analysts with tool that allows them to build effective classification systems without much effort
- **Long-term goal:** An APP that allows anyone to build a classification model from their data (photographs, smart phone data, tweets, etc.)

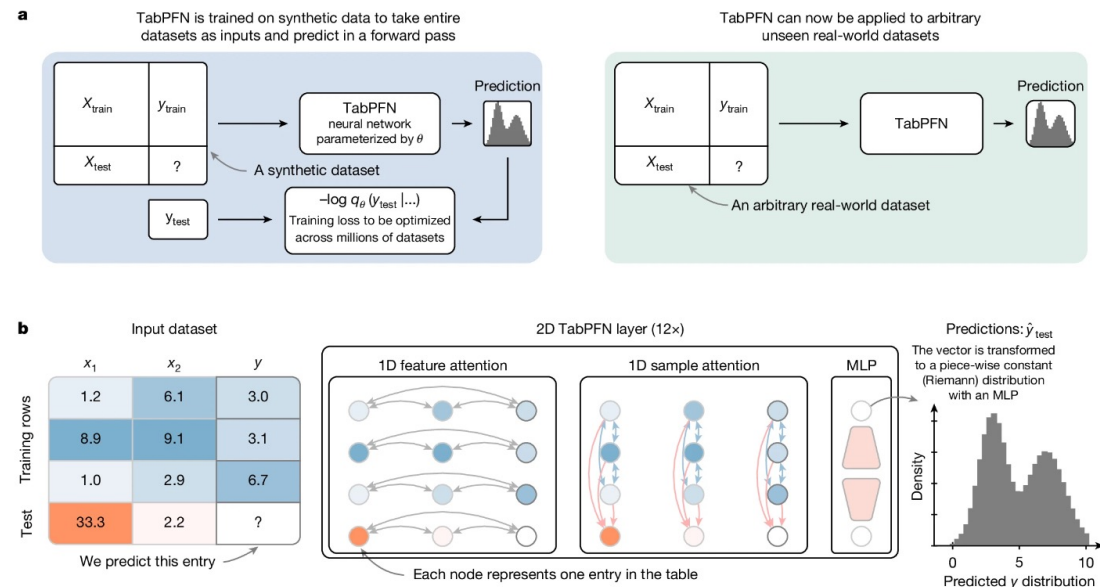


Challenges and research opportunities

- Explainable AutoML models.
- AutoML in feature engineering.
- AutoML for non tabular data.
- Large scale AutoML.
- Transfer learning in AutoML.
- Benchmarking and reproducibility in AutoML.
- Interactive AutoML methods.

Foundational Models of Tabular Data Learning

- **TabPFN (Tabular Prior-Data Fitted Network)** is a transformer-based foundation model designed for fast and accurate classification of tabular datasets.
- Key Features:
 - Ultra-fast predictions in under a second without hyperparameter tuning.
 - Supports both numerical and categorical data.
 - Works effectively for datasets with up to 1,000 samples, 100 features, and 10 classes.
 - Scikit-learn compatible interface for easy integration.



Foundational Models of Tabular Data Learning

- Training Process:**

- TabPFN is trained on millions of synthetic datasets generated from a causal prior.
- This prior reflects realistic patterns found in real-world data.

- Prediction:**

- At inference, the model approximates Bayesian posterior predictions.
- It outputs class probabilities without needing iterative training.

- Architecture:**

- Based on a transformer model adapted for tabular data.

Foundational Models of Tabular Data Learning

•Performance:

- Matches or surpasses gradient boosting and AutoML systems on small datasets.
- Robust even with missing values and feature noise.

•Advantages:

- No hyperparameter tuning required.
- Low computational cost.
- Easy deployment via Python packages or APIs.

Use Cases:

Ideal for quick prototyping, messy structured data tasks.

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.datasets import load_iris
4 from sklearn.model_selection import train_test_split
5 from sklearn.metrics import accuracy_score, classification_report
6 from tabPFN import TabPFNClassifier
7
8 # Load the Iris dataset
9 iris = load_iris()
10 X = iris.data
11 y = iris.target
12
13 # Split the dataset into training and testing sets
14 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
15
16 # Initialize the TabPFN classifier
17 model = TabPFNClassifier(device="cpu") # Change to "cuda" if using a GPU
18
19 # Train the model (TabPFN is designed for quick inference without traditional training)
20 model.fit(X_train, y_train)
21
22 # Make predictions on the test set
23 y_pred = model.predict(X_test)
24
25 # Evaluate the performance
26 accuracy = accuracy_score(y_test, y_pred)
27 report = classification_report(y_test, y_pred, target_names=iris.target_names)
28
29 # Display results
30 print(f"Accuracy: {accuracy:.4f}\n")
```

Foundational Models of Tabular Data Learning

```
1  import numpy as np
2  import pandas as pd
3  from sklearn.datasets import load_iris
4  from sklearn.model_selection import train_test_split
5  from sklearn.metrics import accuracy_score, classification_report
6  from tabPFN import TabPFNClassifier
7
8  # Load the Iris dataset
9  iris = load_iris()
10 X = iris.data
11 y = iris.target
12
13 # Split the dataset into training and testing sets
14 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
15
16 # Initialize the TabPFN classifier
17 model = TabPFNClassifier(device="cpu") # Change to "cuda" if using a GPU
18
19 # Train the model (TabPFN is designed for quick inference without traditional training)
20 model.fit(X_train, y_train)
21
22 # Make predictions on the test set
23 y_pred = model.predict(X_test)
24
25 # Evaluate the performance
26 accuracy = accuracy_score(y_test, y_pred)
27 report = classification_report(y_test, y_pred, target_names=iris.target_names)
28
29 # Display results
30 print(f"Accuracy: {accuracy:.4f}\n")
```